

Company AI Knowledge Base - Page 1

Section 1: Introduction to Generative AI Systems

Generative AI models excel at synthesizing information, generating responses, and retrieving insights across massive collections of unstructured text data.

However, large foundation models have limited context windows and cannot access private enterprise documents without explicit Retrieval-Augmented Generation (RAG).

The first phase of any RAG pipeline is robust document ingestion and chunking.

Company AI Knowledge Base - Page 2

Section 2: Chunking and Embedding Considerations

When splitting large PDF documents, retaining page numbers and headings in metadata is critical for accurate source citations.

Recursive splitting ensures that paragraphs and sentences remain intact rather than being sliced mid-sentence, which preserves semantic coherence.

This completes the reference sample PDF document for testing LangChain loaders.